

```

import os
img_dir = '/tmp/nst'
if not os.path.exists(img_dir):
    os.makedirs(img_dir)
!wget --quiet -P /tmp/nst/ https://upload.wikimedia.org/wikipedia/commons/d/d7/Green_Sea_Turt
!wget --quiet -P /tmp/nst/ https://upload.wikimedia.org/wikipedia/commons/0/0a/The_Great_Wave
!wget --quiet -P /tmp/nst/ https://upload.wikimedia.org/wikipedia/commons/b/b4/Vassily_Kandir
!wget --quiet -P /tmp/nst/ https://upload.wikimedia.org/wikipedia/commons/0/00/Tuebingen_Neck
!wget --quiet -P /tmp/nst/ https://upload.wikimedia.org/wikipedia/commons/6/68/Pillars_of_cre
!wget --quiet -P /tmp/nst/ https://upload.wikimedia.org/wikipedia/commons/thumb/e/ea/Van_Gogh

```

```

import matplotlib.pyplot as plt
import matplotlib as mpl
mpl.rcParams['figure.figsize'] = (10,10)
mpl.rcParams['axes.grid'] = False

```

```

import numpy as np
from PIL import Image
import time
import functools


```

```
import tensorflow as tf
```

```

from tensorflow.python.keras.preprocessing import image as kp_image
from tensorflow.python.keras import models
from tensorflow.python.keras import losses
from tensorflow.python.keras import layers
from tensorflow.python.keras import backend as K

```

 The default version of TensorFlow in Colab will soon switch to TensorFlow 2.x. We recommend you [upgrade](#) now or ensure your notebook will continue to use TensorFlow 1.x via the `%tensorflow_version 1.x` magic: [more info](#).

```

tf.enable_eager_execution()
print("Eager execution: {}".format(tf.executing_eagerly()))

```

 Eager execution: True

```

# Set up some global values here
content_path = '/tmp/nst/Green_Sea_Turtle_grazing_seagrass.jpg'
style_path = '/tmp/nst/The_Great_Wave_off_Kanagawa.jpg'

```

Visualize the input

```
def load_img(path_to_img):
```

```

max_dim = 512
img = Image.open(path_to_img)

long = max(img.size)
scale = max_dim/long
img = img.resize((round(img.size[0]*scale), round(img.size[1]*scale)), Image.ANTIALIAS)
img = kp_image.img_to_array(img)

# We need to broadcast the image array such that it has a batch dimension
img = np.expand_dims(img, axis=0)
return img

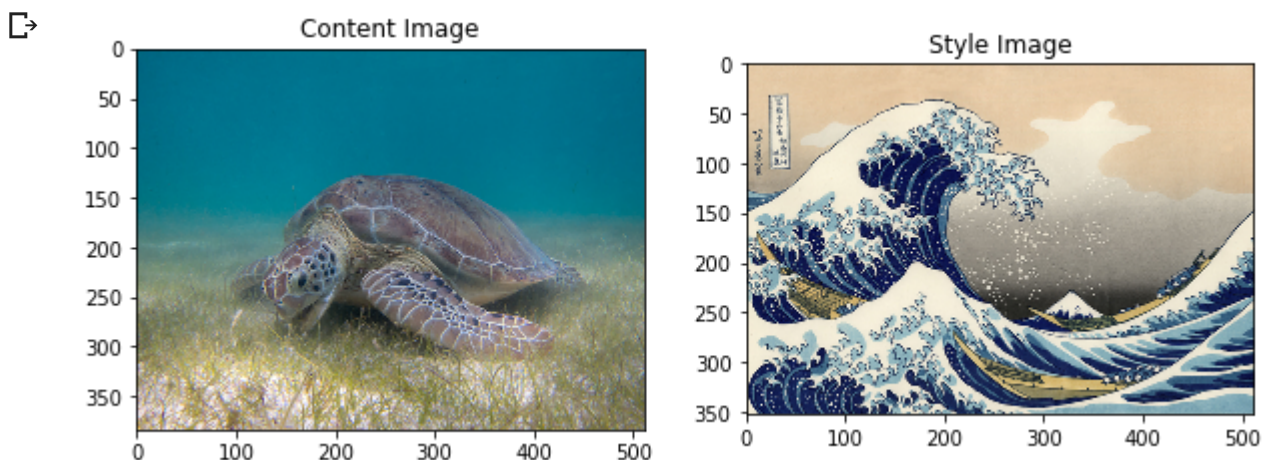
def imshow(img, title=None):
    # Remove the batch dimension
    out = np.squeeze(img, axis=0)
    # Normalize for display
    out = out.astype('uint8')
    plt.imshow(out)
    if title is not None:
        plt.title(title)
    plt.imshow(out)

plt.figure(figsize=(10,10))

content = load_img(content_path).astype('uint8')
style = load_img(style_path).astype('uint8')
plt.subplot(1, 2, 1)
imshow(content, 'Content Image')

plt.subplot(1, 2, 2)
imshow(style, 'Style Image')
plt.show()

```



Prepare the data

```
def load_and_process_img(path_to_img):
    img = load_img(path_to_img)
    img = tf.keras.applications.vgg19.preprocess_input(img)
    return img

def deprocess_img(processed_img):
    x = processed_img.copy()
    if len(x.shape) == 4:
        x = np.squeeze(x, 0)
    assert len(x.shape) == 3, ("Input to deprocess image must be an image of "
                               "dimension [1, height, width, channel] or [height, width, channel]")
    if len(x.shape) != 3:
        raise ValueError("Invalid input to deprocessing image")

    # perform the inverse of the preprocessing step
    x[:, :, 0] += 103.939
    x[:, :, 1] += 116.779
    x[:, :, 2] += 123.68
    x = x[:, :, ::-1]

    x = np.clip(x, 0, 255).astype('uint8')
    return x

# Content layer where will pull our feature maps
content_layers = ['block5_conv2']

# Style layer we are interested in
style_layers = ['block1_conv1',
                'block2_conv1',
                'block3_conv1',
                'block4_conv1',
                'block5_conv1'
                ]

num_content_layers = len(content_layers)
num_style_layers = len(style_layers)
```

Build the Model

```
def get_model():
    """ Creates our model with access to intermediate layers.
```

This function will load the VGG19 model and access the intermediate layers. These layers will then be used to create a new model that will take input image and return the outputs from these intermediate layers from the VGG model.

Returns:

returns a keras model that takes image inputs and outputs the style and content intermediate layers.

"""

```
# Load our model. We load pretrained VGG, trained on imagenet data
vgg = tf.keras.applications.vgg19.VGG19(include_top=False, weights='imagenet')
vgg.trainable = False
# Get output layers corresponding to style and content layers
style_outputs = [vgg.get_layer(name).output for name in style_layers]
content_outputs = [vgg.get_layer(name).output for name in content_layers]

model_outputs = style_outputs + content_outputs
# Build model
return models.Model(vgg.input, model_outputs)
```

Computing content loss

```
def get_content_loss(base_content, target):
    return tf.reduce_mean(tf.square(base_content - target))
```

Computing Style Loss

```
def gram_matrix(input_tensor):
    # We make the image channels first
    channels = int(input_tensor.shape[-1])
    a = tf.reshape(input_tensor, [-1, channels])
    n = tf.shape(a)[0]
    gram = tf.matmul(a, a, transpose_a=True)
    return gram / tf.cast(n, tf.float32)

def get_style_loss(base_style, gram_target):
    """Expects two images of dimension h, w, c"""
    # height, width, num filters of each layer
    # We scale the loss at a given layer by the size of the feature map and the number of filters
    height, width, channels = base_style.get_shape().as_list()
    gram_style = gram_matrix(base_style)

    return tf.reduce_mean(tf.square(gram_style - gram_target)) / (4. * (channels ** 2) * (width
```

```
def get_feature_representations(model, content_path, style_path):
    """Helper function to compute our content and style feature representations.

    This function will simply load and preprocess both the content and style
    images from their path. Then it will feed them through the network to obtain
    the outputs of the intermediate layers.

    Arguments:
        model: The model that we are using.
        content_path: The path to the content image.
        style_path: The path to the style image

    Returns:
        returns the style features and the content features.
    """
    # Load our images in
    content_image = load_and_process_img(content_path)
    style_image = load_and_process_img(style_path)

    # batch compute content and style features
    style_outputs = model(style_image)
    content_outputs = model(content_image)

    # Get the style and content feature representations from our model
    style_features = [style_layer[0] for style_layer in style_outputs[:num_style_layers]]
    content_features = [content_layer[0] for content_layer in content_outputs[num_style_layers:]]
    return style_features, content_features
```

Computing the loss and gradients

```
def compute_loss(model, loss_weights, init_image, gram_style_features, content_features):
    """This function will compute the loss total loss.

    Arguments:
        model: The model that will give us access to the intermediate layers
        loss_weights: The weights of each contribution of each loss function.
            (style weight, content weight, and total variation weight)
        init_image: Our initial base image. This image is what we are updating with
            our optimization process. We apply the gradients wrt the loss we are
            calculating to this image.
        gram_style_features: Precomputed gram matrices corresponding to the
            defined style layers of interest.
        content_features: Precomputed outputs from defined content layers of
            interest.

    Returns:
        returns the total loss, style loss, content loss, and total variational loss
```

```

"""
style_weight, content_weight = loss_weights

# Feed our init image through our model. This will give us the content and
# style representations at our desired layers. Since we're using eager
# our model is callable just like any other function!
model_outputs = model(init_image)

style_output_features = model_outputs[:num_style_layers]
content_output_features = model_outputs[num_style_layers:]

style_score = 0
content_score = 0

# Accumulate style losses from all layers
# Here, we equally weight each contribution of each loss layer
weight_per_style_layer = 1.0 / float(num_style_layers)
for target_style, comb_style in zip(gram_style_features, style_output_features):
    style_score += weight_per_style_layer * get_style_loss(comb_style[0], target_style)

# Accumulate content losses from all layers
weight_per_content_layer = 1.0 / float(num_content_layers)
for target_content, comb_content in zip(content_features, content_output_features):
    content_score += weight_per_content_layer * get_content_loss(comb_content[0], target_content)

style_score *= style_weight
content_score *= content_weight

# Get total loss
loss = style_score + content_score
return loss, style_score, content_score

def compute_grads(cfg):
    with tf.GradientTape() as tape:
        all_loss = compute_loss(**cfg)
    # Compute gradients wrt input image
    total_loss = all_loss[0]
    return tape.gradient(total_loss, cfg['init_image']), all_loss

```

Optimization loop

```

import IPython.display

def run_style_transfer(content_path,
                      style_path,
                      num_iterations=1000,
                      content_weight=1e3,

```

```

style_weight = 1e-2):

display_num = 100
# We don't need to (or want to) train any layers of our model, so we set their trainability
# to false.
model = get_model()
for layer in model.layers:
    layer.trainable = False

# Get the style and content feature representations (from our specified intermediate layers)
style_features, content_features = get_feature_representations(model, content_path, style_path)
gram_style_features = [gram_matrix(style_feature) for style_feature in style_features]

# Set initial image
init_image = load_and_process_img(content_path)
init_image = tf.Variable(init_image, dtype=tf.float32)
# Create our optimizer
opt = tf.train.AdamOptimizer(learning_rate=10.0)

# For displaying intermediate images
iter_count = 1

# Store our best result
best_loss, best_img = float('inf'), None

# Create a nice config
loss_weights = (style_weight, content_weight)
cfg = {
    'model': model,
    'loss_weights': loss_weights,
    'init_image': init_image,
    'gram_style_features': gram_style_features,
    'content_features': content_features
}

# For displaying
plt.figure(figsize=(15, 15))
num_rows = (num_iterations / display_num) // 5
start_time = time.time()
global_start = time.time()

norm_means = np.array([103.939, 116.779, 123.68])
min_vals = -norm_means
max_vals = 255 - norm_means

imgs = []

for i in range(num_iterations):
    grads, all_loss = compute_grads(cfg)
    loss, style_score, content_score = all_loss
    #grads, _ = tf.clip_by_global_norm(grads, 5.0)
    opt.apply_gradients([(grads, init_image)])
    clipped = tf.clip_by_value(init_image, min_vals, max_vals)

```

```

init_image.assign(clipped)
end_time = time.time()

if loss < best_loss:
    # Update best loss and best image from total loss.
    best_loss = loss
    best_img = init_image

if i % display_num == 0:

    # Use the .numpy() method to get the concrete numpy array
    plot_img = init_image.numpy()
    plot_img = deprocess_img(plot_img)
    imgs.append(plot_img)
    IPython.display.clear_output(wait=True)
    IPython.display.display_png(Image.fromarray(plot_img))
    print('Iteration: {}'.format(i))
    print('Total loss: {:.4e}, '
          'style loss: {:.4e}, '
          'content loss: {:.4e}, '
          'time: {:.4f}s'.format(loss, style_score, content_score, time.time() - start_time))
    start_time = time.time()

    # Display intermediate images
    if iter_count > num_rows * 5: continue
    plt.subplot(num_rows, 5, iter_count)
    # Use the .numpy() method to get the concrete numpy array
    plot_img = init_image.numpy()
    plot_img = deprocess_img(plot_img)
    plt.imshow(plot_img)
    plt.title('Iteration {}'.format(i + 1))

    iter_count += 1
print('Total time: {:.4f}s'.format(time.time() - global_start))

return best_img, best_loss

best, best_loss = run_style_transfer(content_path,
                                     style_path, num_iterations=1001)

```

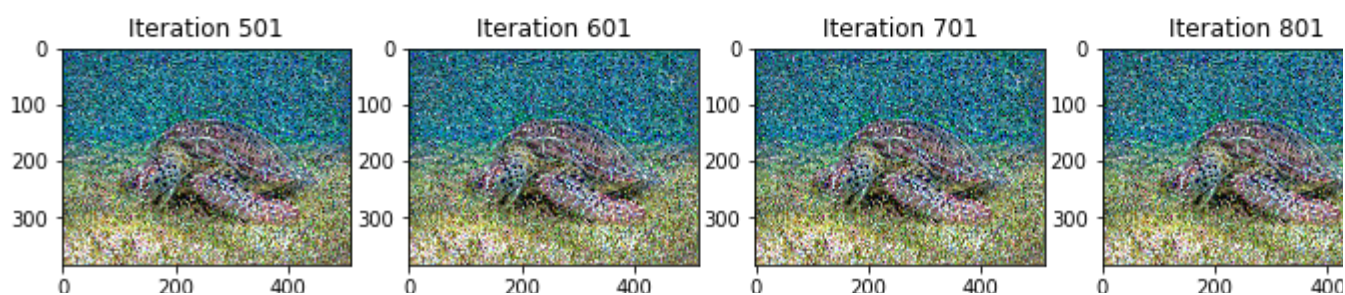
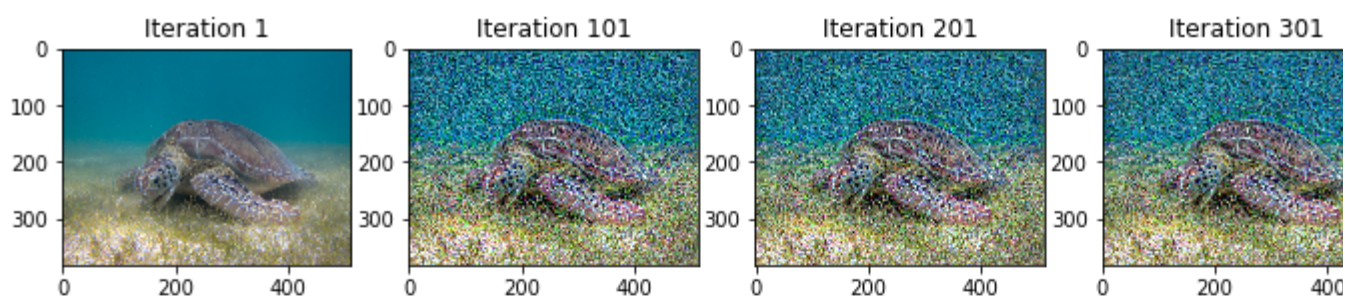




Iteration: 1000

Total loss: 6.6705e+03, style loss: 3.9069e-05, content loss: 6.6705e+03, time: 7.8961s

Total time: 78.9736s



```
def show_results(best_image, content_path, style_path, show_large_final=True):

    plt.figure(figsize=(10, 5))
    content = load_img(content_path)
    style = load_img(style_path)

    plt.subplot(1, 2, 1)
    imshow(content, 'Content Image')

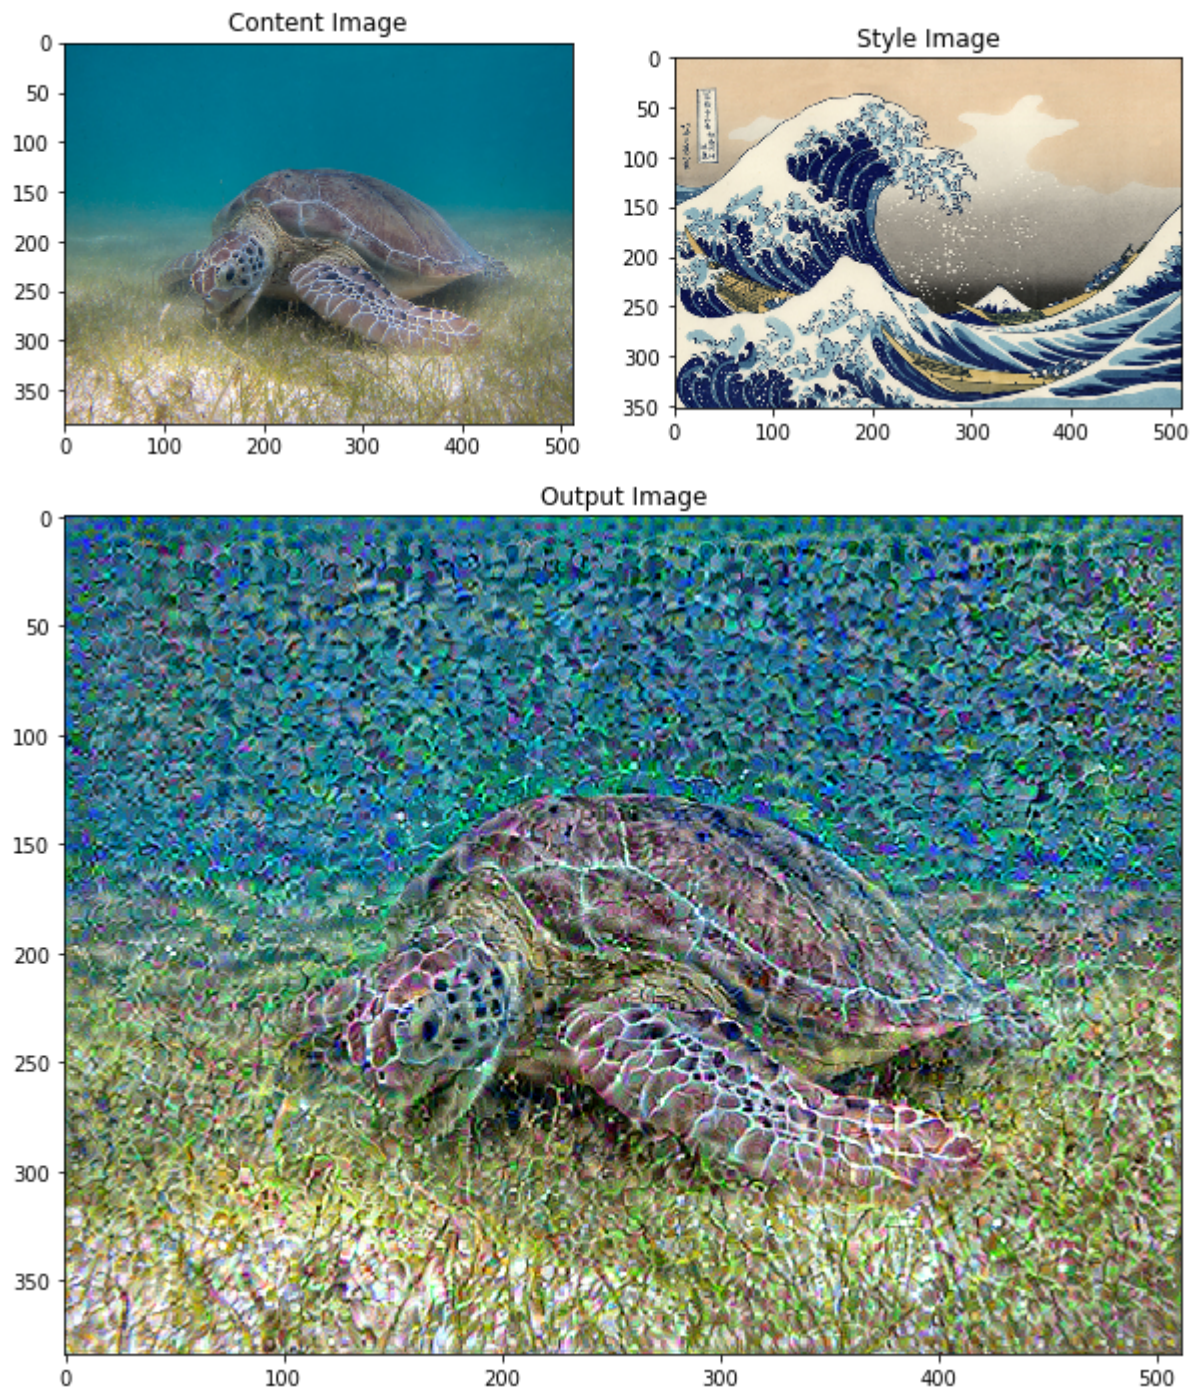
    plt.subplot(1, 2, 2)
    imshow(style, 'Style Image')

    if show_large_final:
        plt.figure(figsize=(10, 10))

        best_image = best_image.numpy()
        best_image = deprocess_img(best_image)
        plt.imshow(best_image)
        plt.title('Output Image')
        plt.show()

show_results(best, content_path, style_path)
```





```
plt.imshow(deprocess_img(best))
```

